

# Cloud Applications Architecture



Course 1 - Intro (Course & Cloud)

# Why Would You Be Interested?

Cloud is similar to the whole IT industry - relevant for all other industries.

Consistently in top 3 most demanded (hard - as opposed to soft) skills.

It will be useful no matter what you choose to do after.

We will have fun.

# What Does CAA Mean?

Designing and building applications (IT systems) which leverage the possibilities (services) offered by cloud environments to achieve greater:

- Scalability
- Security
- Cost Efficiency
- Availability
- Productivity

In other words, build systems we can trust and have a good time while doing it.

# General Info

## Course

- Focused on theoretical concepts, case studies, and demos
- Probably 50% of final grade
- Involvement is worth more than learning for the exam
- Most likely quiz exam using Moodle (maybe some oral questions?)

## Lab

- Hands-on with AWS
- Migrate an app from **on-premises** to **cloud**
- 50% of final grade
- Architecture design quiz

# General Info

- Topics Overview

# What Is the Traditional Approach?

- For individuals and small companies: develop a monolithic app, contact a hosting company, rent a server, install and configure relevant tools, copy the app with FTP, run the app.
- For large corporations: Discuss the budget, approve the budget, build a data center, go over the budget, hire people to manage the data center, deploy apps, upgrade data center.

# Cloud Intro

# What does Cloud mean?

In simple terms, cloud computing means on-demand availability of various computing resources. Typical benefits of using the cloud include:

- Flexibility
- Cost
- Speed
- Security
- Performance
- Reliability
- ...



# History

First public offering came from a large company (Amazon) as a result of over-provisioning (to handle holiday sales) in 2006.

In 2008, Google launched App Engine, one of the first public PaaS solutions.

# Main Providers



# Gartner Magic Quadrant

Figure 1. Magic Quadrant for Cloud Infrastructure and Platform Services



Figure 1. Magic Quadrant for Cloud AI Developer Services



# Cloud Classification

## Deployment Models:

- **Public**
- **Private**
- **Community**
- **Hybrid**

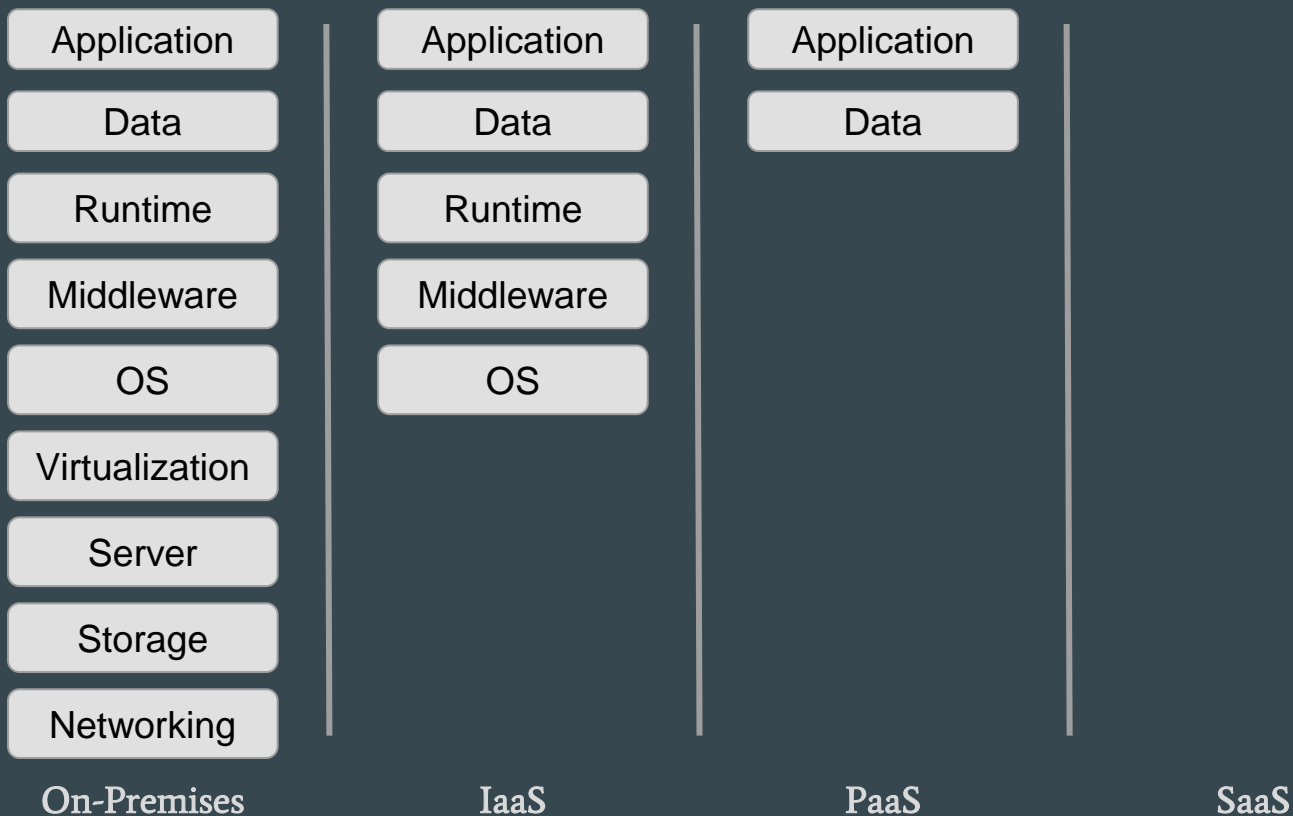
## Service Models:

- **IaaS (infrastructure as a service)**
- **PaaS (platform as a service)**
- **SaaS (software as a service)**

Many \*aaS nowadays: BaaS, DBaaS, FaaS (most of them are PaaS)

[The NIST Model](#)

# IaaS, PaaS, SaaS - What We Manage



# Examples for Each Service Model

## IaaS

Usually virtual machine services. Other service models are built on top of them.

E.g. Amazon EC2, Google's GCE etc.

## PaaS

Services that let you develop your solution without worrying about how the system runs.

E.g. Cloudfoundry, Google App Engine, Azure App Service, managed databases (RDS, CloudSQL etc.)

## SaaS

No implementation, you just use the provided functionality. Interact with them through UI or API.

E.g. Google Suite, Office365, Stripe, Dropbox, Mailgun

# Working with the Cloud

1. Create an account
2. Setup billing
  - a. Most providers offer decent free tiers
3. Configure and launch resources
  - a. From the web console
  - b. Using the CLI
  - c. Infrastructure as Code
4. \*Stop/terminate them when not used
  - a. Usually billing stops when stopping the resources

# Standards/Compliance

Assure customers best practices are followed

- [CSA](#) (Cloud Security Alliance) - Best practices for cloud security
- [ISO 9001](#) - Global Quality Standard
- [ISO 27001](#) - Information Security Management
- [ISO 27017](#) - 27001 in the context of cloud
- [ISO 27018](#) - Personal data protection in cloud
- [PCI DSS](#) - Card payments security
- [HIPAA](#) - Health information protection

Running on complaint infrastructure makes our products also compliant



# Pricing Models

Most services are billed based on the actual usage:

- **Time based:** pay per second/hour etc of the service being online
  - Usually applies to IaaS and PaaS offerings
- **Usage based:** pay per processing unit (request, transaction, read, write etc)
  - Mostly for PaaS/SaaS

Prices can be **fixed**, **tiered** and/or **dynamic**.

There is also the **subscription** based model: agree to pay each month a certain price and have access to the agreed services. Usually applies to SaaS offerings and some providers - e.g. SAP offers both

# Pricing Models - Time vs Usage

## Time Based

Usually applies to virtual machines, (relational) databases and general PaaS services (e.g. Cloudfoundry).

You pay the same for any amount of traffic.

Easy to control costs.

## Usage Based

Usually applies to fully managed services (FaaS, DBaaS - usually NoSQL) and storage services.

You pay based on the traffic.

You cannot throttle the usage without stopping the service.

Some services offer both - e.g.

[Azure API Management](#)

# Pricing Models - Fixed, Tiered, Dynamic

## Fixed

You pay the same no matter what. Usually applies to on-demand virtual machines

Instance	vCPU(s)	RAM	Temporary storage	Pay as you go
B1S	1	1 GiB	4 GiB	€0.0135/hour
B1MS	1	2 GiB	4 GiB	€0.0234/hour
B2S	2	4 GiB	8 GiB	€0.0473/hour

[Azure Windows Virtual Machines Pricing](#)

## Tiered

You pay less per unit the more you use it. Usually applies to storage services and bandwidth cost.

Storage pricing	
S3 Standard - General purpose storage for any type of data, typically used for frequently accessed data	
First 50 TB / Month	\$0.023 per GB
Next 450 TB / Month	\$0.022 per GB
Over 500 TB / Month	\$0.021 per GB

## Dynamic

Usually based on supply and demand. More spare capacity, smaller prices.

Linux/UNIX Usage		Windows Usage
General Purpose - Current Generation		
a1.medium	\$0.0084 per Hour	N/A*
a1.large	\$0.0234 per Hour	N/A*
a1.xlarge	\$0.0336 per Hour	N/A*
a1.2xlarge	\$0.0672 per Hour	N/A*
a1.4xlarge	\$0.1343 per Hour	N/A*
a1.metal	\$0.13	
t2.micro	\$0.00	
t2.small	\$0.00	
		Linux/UNIX Usage
		Windows Usage
General Purpose - Current Generation		
a1.medium	\$0.005 per Hour	N/A*
a1.large	\$0.0098 per Hour	N/A*
a1.xlarge	\$0.0197 per Hour	N/A*
a1.2xlarge	\$0.0394 per Hour	N/A*
a1.4xlarge	\$0.0788 per Hour	N/A*
a1.metal	\$0.0788 per Hour	N/A*
t2.micro	\$0.0035 per Hour	\$0.0081 per Hour
t2.small	\$0.0069 per Hour	\$0.0159 per Hour

# Pricing Models - Discounts

## Credits

Most providers offer **sign-up** credits, **educational** credits and **start-up** credits. You can also earn them from various contests or partnerships.

They also provide always free service tiers.

## Reservation

If you know you will be using a certain service for a long time, you can commit to it and get a considerable price cut.

E.g. AWS Reserved Instances and Savings Plans can reduce the cost by 30 to 70%

## Sustained Use

Some providers will automatically apply discounts to your bill if you use certain resources continuously.

E.g. Google Cloud sustained use discounts for GCE and App Engine (30% if the instance is running the entire month)

# Guidelines & Principles

# Architecture Frameworks

**AWS Well-Architected**

**Google Architecture Framework**

**Azure Well-Architected**

**Operational Excellence**

- Automation, Monitoring, Testing

**Security**

- IAM, Defence in depth, Traceability

**Reliability**

- Monitoring, Scaling, Testing

**Performance**

- Scaling, Choosing right, Design

**Cost Optimization**

- Choosing right, scaling, Understanding the services and offers

# The 12-Factor App [Link](#)

Codebase

Port Binding

Dependencies

Concurrency

Config

Disposability

Backing Services

Dev/Prod Parity

Build, release, run

Logs

Processes

Admin Processes

[Pragmatic video about it](#)

# Relevant Links

1. [List with engineering blogs from top companies](#) (most relevant: Netflix, Uber, Atlassian)
2. [AWS re:Invent 2019 Playlist](#)
3. [Google Cloud Next `19 Playlist](#)
4. [Microsoft Developer Youtube Channel](#)
5. [Last Week in AWS](#)



# Summary

1. Why you would care
2. Involvement is important
3. The cloud and how it compares to on-premises
4. Providers
5. Cloud classification
6. Standards
7. Pricing models
8. Principles